


# Matrix multiplication

Serge Kruk

October 5, 2021

# Matrix multiplication

$$A_{p \times q} \times B_{q \times r} = C_{p \times r}$$



# Design an algorithm to perform matrix multiplication

By the method you learned in Linear Algebra.

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \cdot \begin{bmatrix} 9 & 8 \\ 7 & 6 \\ 5 & 4 \end{bmatrix} = \begin{bmatrix} 9 + 14 + 15 & 8 + 12 + 12 \\ 36 + 35 + 30 & 32 + 30 + 24 \end{bmatrix}$$
$$= \begin{bmatrix} 38 & 32 \\ 101 & 86 \end{bmatrix}$$

# Design an algorithm to perform matrix multiplication

- Loop over every row of the first matrix.
  - Loop over every column of the second matrix.
    - Do an dot product of the selected row/column pair.

# Design an algorithm to perform matrix multiplication

```
1 def size(A):
2     return len(A), len(A[0])
3
4 def mm(A,B):
5     (p,q) = size(A)
6     (q,r) = size(B)
7     C = [[0 for _ in range(r)] for _ in range(p)]
8     for i in range(p):
9         for j in range(q):
10            for k in range(r):
11                C[i][k] += A[i][j] * B[j][k]
12     return C
```

# Little test

```
A=[[1,2,3],[4,5,6]]  
B=[[9,8],[7,6],[5,4]]  
mm(A,B)
```

```
38 32  
101 86
```

Let us count multiplications and additions

$$T(p, q, r) = \sum_p \sum_r \sum_q 1$$

- $\Theta(pqr)$
- $\Theta(n^3)$  if all matrices are size  $n \times n$ 
  - $n^3$  Multiplications
  - $n^3$  Additions

# We could reduce the number of additions a bit

- The innermost operation is an inner product

```
1 def mmf(A,B):
2     (p,q) = size(A)
3     (q,r) = size(B)
4     C = [[0 for _ in range(r)] for _ in range(p)]
5     for i in range(p):
6         for j in range(r):
7             C[i][j] = sum(A[i][k] * B[k][j] for k in range(q))
8     return C
```

A=[[1,2,3],[4,5,6]]

B=[[9,8],[7,6],[5,4]]

mmf(A,B)

38 32  
101 86



# We could reduce the number of additions a bit

- We avoid one addition per inner product.
- There are  $n^2$  inner products on an  $n \times n$  matrix.
- Hence the counts are now
  - $n^3$  multiplications
  - $n^3 - n^2$  additions
- Don't lose sleep over this. It is the kind of micro-optimization that makes or breaks numerical code but is not important for asymptotic analysis.

Note also, in passing ...

```
for i in range(p):  
    for j in range(q):  
        for k in range(r):  
            C[i][k] += A[i][j] * B[j][k]
```

or

```
for k in range(r):  
    for i in range(p):  
        for j in range(q):  
            C[i][k] += A[i][j] * B[j][k]
```

or

```
for i in range(p):  
    for k in range(r):  
        for j in range(q):  
            C[i][k] += A[i][j] * B[j][k]
```

or ...

# How are matrices stored in computer memory?

- Say

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

Then, in memory it could be either of these

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 |
|---|---|---|---|---|---|

**Row major (C family, Java, Octave)**

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 1 | 4 | 2 | 5 | 3 | 6 |
|---|---|---|---|---|---|

**Column major (FORTRAN, Matlab, Julia)**

# Data access pattern

| Loop order | Data access              |
|------------|--------------------------|
| ijk        | A by row, B by column    |
| jik        | A by row, B by column    |
| ikj        | B by row, C by row       |
| jki        | A by column, C by column |
| kij        | B by row, C by row       |
| kji        | A by column, C by column |

- Again, do not lose sleep over this. I mention it to highlight the kinds of nightmares numerical analysts face daily.
- For more on this, the bible is **Matrix computations** by Golub and Van Loan.

# Back to our regular programme

Now let's go back to our problem: how to multiply matrices.

# Challenge: do better than $n^3$

- Do you think it is possible?
- We will need to look at cost in more detail.

# Assumption

For a while, let us assume that  $n$  is a power of 2.

# Demo of block multiplication

```
octave:2> A11=round(10*rand(2,2))
```

```
A11 =
```

```
5    3
```

```
0   10
```

```
octave:3> A12=round(10*rand(2,2))
```

```
A12 =
```

```
7    8
```

```
0    0
```

```
octave:4> A21=round(10*rand(2,2))
```

```
A21 =
```

```
0    3
```

```
2    6
```

```
octave:5> A22=round(10*rand(2,2))
```

```
A22 =
```

```
2    8
```

```
2    8
```



# Demo of block multiplication

```
octave:6> B11=round(10*rand(2,2))
```

```
B11 =
```

```
9    10
```

```
3     6
```

```
octave:7> B12=round(10*rand(2,2))
```

```
B12 =
```

```
7     9
```

```
5     9
```

```
octave:8> B21=round(10*rand(2,2))
```

```
B21 =
```

```
8     9
```

```
3     8
```

```
octave:9> B22=round(10*rand(2,2))
```

```
B22 =
```

```
3     7
```

```
6     5
```

# Demo of block multiplication

```
octave:10> A=[A11,A12;A21,A22]
```

A =

|   |    |   |   |
|---|----|---|---|
| 5 | 3  | 7 | 8 |
| 0 | 10 | 0 | 0 |
| 0 | 3  | 2 | 8 |
| 2 | 6  | 2 | 8 |

```
octave:11> B=[B11,B12;B21,B22]
```

B =

|   |    |   |   |
|---|----|---|---|
| 9 | 10 | 7 | 9 |
| 3 | 6  | 5 | 9 |
| 8 | 9  | 3 | 7 |
| 3 | 8  | 6 | 5 |

# Demo of block multiplication

```
octave:12> A*B
```

```
ans =
```

|     |     |     |     |
|-----|-----|-----|-----|
| 134 | 195 | 119 | 161 |
| 30  | 60  | 50  | 90  |
| 49  | 100 | 69  | 81  |
| 76  | 138 | 98  | 126 |

```
octave:16> [A11*B11+A12*B21, A11*B12+A12*B22;  
            A21*B11+A22*B21, A21*B12+A22*B22]
```

```
ans =
```

|     |     |     |     |
|-----|-----|-----|-----|
| 134 | 195 | 119 | 161 |
| 30  | 60  | 50  | 90  |
| 49  | 100 | 69  | 81  |
| 76  | 138 | 98  | 126 |

Consider a  $2m \times 2m$  block matrix

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

Where

$$C_{11} = A_{11}B_{11} + A_{12}B_{21}$$

$$C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21}$$

$$C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

- 8 multiplications (of  $m \times m$  matrices)
- 4 additions (of  $m \times m$  matrices)

# Assuming we do the submatrix operations in standard way

- $A_{ij}B_{kl}$ 
  - $(m-1)m^2 = m^3 - m^2$  additions
  - $m^3$  multiplications
- $A_{ij}B_{kl} + A_{op}B_{qp}$ 
  - $2(m^3 - m^2) + m^2$  additions
  - $2m^3$  multiplications
- For the four blocks
  - $4(2m^3 - m^2) = 8m^3 - 4m^2$  additions
  - $8m^3$  multiplications

# Strassen's first brilliant idea

He claims that

$$P_1 = (A_{11} + A_{22})(B_{11} + B_{22})$$

$$P_2 = (A_{21} + A_{22})B_{11}$$

$$P_3 = A_{11}(B_{12} - B_{22})$$

$$P_4 = A_{22}(B_{21} - B_{11})$$

$$P_5 = (A_{11} + A_{12})B_{22}$$

$$P_6 = (A_{21} - A_{11})(B_{11} + B_{12})$$

$$P_7 = (A_{12} - A_{22})(B_{21} + B_{22})$$

$$C_{11} = P_1 + P_4 - P_5 + P_7$$

$$C_{12} = P_3 + P_5$$

$$C_{21} = P_2 + P_4$$

$$C_{22} = P_1 - P_2 + P_3 + P_6$$

# Strassen's first brilliant idea

Does it work?

```
octave:26> P1=(A11+A22)*(B11+B22)
```

```
P1 =
```

```
183    240
```

```
186    232
```

```
octave:20> P4=A22*(B21-B11)
```

```
P4 =
```

```
-2     14
```

```
-2     14
```

```
octave:21> P5=(A11+A12)*B22
```

```
P5 =
```

```
102    139
```

```
60     50
```

```
octave:23> P7=(A12-A22)*(B21+B22)
```

```
P7 =
```

```
55     80
```

```
-94    -136
```

# Strassen's first brilliant idea

```
octave:27> P1+P4-P5+P7
```

```
ans =
```

```
134    195
```

```
30     60
```

```
octave:28> A*B
```

```
ans =
```

```
134    195    119    161
```

```
30     60     50     90
```

```
49    100     69     81
```

```
76    138     98    126
```



# Strassen's first brilliant idea

- How would you prove that it works?
- By showing that the specified sums yield the product.

# Strassen's first brilliant idea (runtime)

Count multiplications and additions.

$$P_1 = (A_{11} + A_{22})(B_{11} + B_{22})$$

$$P_2 = (A_{21} + A_{22})B_{11}$$

$$P_3 = A_{11}(B_{12} - B_{22})$$

$$P_4 = A_{22}(B_{21} - B_{11})$$

$$P_5 = (A_{11} + A_{12})B_{22}$$

$$P_6 = (A_{21} - A_{11})(B_{11} + B_{12})$$

$$P_7 = (A_{12} - A_{22})(B_{21} + B_{22})$$

$$C_{11} = P_1 + P_4 - P_5 + P_7$$

$$C_{12} = P_3 + P_5$$

$$C_{21} = P_2 + P_4$$

$$C_{22} = P_1 - P_2 + P_3 + P_6$$

# Strassen's first brilliant idea (runtime)

- We have 7 matrix multiplications and 18 additions.
- Doing this, using standard multiplication at the block level.
  - $7(m^3 - m^2) + 18m^2 = 7m^3 + 11m^2$  scalar additions
  - $7m^3$  scalar multiplications

## Second surprise

By simple algebra, we shave some operations

- By block view
  - $8m^3 - 4m^2$  additions
  - $8m^3$  multiplications
- By Strassen's first idea and block view
  - $7m^3 + 11m^2$  additions
  - $7m^3$  multiplications

If  $m$  is large enough Strassen's wins!

- Assuming multiplication is at least as costly as addition.
- And that the data access pattern doesn't kill cache access.
- Assuming the pipeline is filled properly
- Etc. . .

# Strassen's second brilliant idea

Why stop at one level, recurse! Down to a certain small size.

# Strassen's runtime

Assumption (made to simplify analysis)

- $n = 2^q$
- $n_0 = 2^d$

What is the proper recurrence?

# Strassen's runtime

Assumption (made to simplify analysis)

- $n = 2^q$
- $n_0 = 2^d$

$$T(n) = \begin{cases} n_0^3 & n \leq n_0 \\ 7T(n/2) + c_0 n^2 & \text{otherwise} \end{cases}$$

# Strassen's runtime

$$\begin{aligned}T(n) &= 7T(n/2) + c_0 n^2 \\&= 7[7T(n/4) + c_{11}(n/2)^2] + c_0 n^2 &= 7^2 T(n/4) + c_1 n^2 \\&= 7^2[7T(n/8) + c_{22}(n/4)^2] + c_1 n^2 &= 7^3 T(n/8) + c_2 n^2 \\&= \dots \\&= 7^{\log_2 n} n_0^3 + c_{\log_2 n} n^2 \\&= \Theta(7^{\log_2 n})\end{aligned}$$



By the master theorem, you get  $\Theta(n^{\log_2 7})$

- Equivalent to our bound (maybe you want to prove that)
- $\Theta(7^{\log_2 n}) = \Theta(n^{\log_2 7}) \approx \Theta(n^{2.8})$

- Assume the matrix is  $m \times m$ 
  - If  $m$  is even then we simply divide in half
  - If  $m$  is odd, we pad with a row/column of zeros
- We stop the recursion when  $m$  is smallish
  - Nowadays that may mean in the low thousands depending on hardware/software/cache

# Strassen's algorithm

**WARNING:** This is pseudo-code. It does not compile.

```
def strassen(A,B,n,n0=1024):
    if n <= n0:
        return mm(A,B)
    else:
        m,u,v = n//2,range(m),range(m+1,n)
        P1 = strassen(A(u,u)+A(v,v),B(u,u)+B(v,v),m,n0)
        P2 = strassen(A(v,u)+A(v,v),B(u,u),m,n0)
        P3 = strassen(A(u,u), B(u,v)-B(v,v),m,n0)
        ...
        C(u,u) = P1 + P4 - P5 + P7
        C(u,v) = P3 + P5
        C(v,u) = P2 + P4
        C(v,v) = P1 - P2 + P3 + P6
    return C
```

# Strassen in Octave (This is executable, if slooooooow)

```
function [C] = strassen(A, B, n0=3)
    [row,common] = size(A);
    [common0, column] = size(B);
    assert(common==common0);
    if row <= n0 || column <= n0,
        C = A*B;
    else
        if rem(row, 2) == 1,
            A = [A; zeros(1, common)];
        endif;
        if rem(column, 2) == 1,
            B = [B zeros(common, 1)];
        endif;
        if rem(common, 2) == 1,
            A = [A zeros(size(A, 1), 1)];
            B = [B; zeros(1, size(B, 2))];
        endif;
```

# Strassen in Octave

```
[m,o] = size(A);  
[o,n] = size(B);  
A11 = A(1:m/2, 1:n/2 );  
A12 = A(1:m/2, n/2+1: n);  
A21 = A(m/2+1 :m, 1:n/2);  
A22 = A(m/2+1 :m ,n/2+1: n);  
B11 = B(1:n/2, 1:o/2 );  
B12 = B(1:n/2, o/2+1: o);  
B21 = B(n/2+1 :n, 1:o/2);  
B22 = B(n/2+1 :n, o/2+1: o);  
P1 = strassen(A11 + A22, B11 + B22, n0);  
P2 = strassen(A21 + A22, B11, n0);  
P3 = strassen(A11, B12 - B22, n0);  
P4 = strassen(A22, B21 - B11, n0);  
P5 = strassen(A11 + A12, B22, n0);  
P6 = strassen(A21 - A11, B11 + B12, n0);  
P7 = strassen(A12 - A22, B21 + B22, n0);
```

# Strassen in Octave

```
C = [(P1 + P4 - P5 + P7) (P3 + P5);  
      (P2 + P4) (P1 + P3 - P2 + P6)](1:row, 1:column);  
endif;  
endfunction;
```

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 256 | 259 | 256 | 329 | 212 | 293 |
| 114 | 117 | 172 | 166 | 128 | 141 |
| 143 | 102 | 139 | 220 | 115 | 191 |
| 156 | 175 | 109 | 174 | 105 | 175 |
| 115 | 145 | 80  | 140 | 92  | 160 |
| 149 | 127 | 197 | 224 | 161 | 200 |

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.



# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!
- What is the optimal runtime?

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!
- What is the optimal runtime?
  - Nobody knows.

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!
- What is the optimal runtime?
  - Nobody knows.
  - Currently theoretical best is  $\Theta(n^{2.3728639})$  (2014, François Le Gall)

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!
- What is the optimal runtime?
  - Nobody knows.
  - Currently theoretical best is  $\Theta(n^{2.3728639})$  (2014, François Le Gall)
  - Currently Strassen's is better **in practice**.

# Theoretical conclusions

- Can we do better than  $\Theta(n^3)$ ?
  - Yes! The first was Strassen  $\Theta(n^{2.8})$ , but many followed.
- Can we do better than  $\Theta(n^2)$ ?
  - No! We need to output  $n^2$  elements!
- What is the optimal runtime?
  - Nobody knows.
  - Currently theoretical best is  $\Theta(n^{2.3728639})$  (2014, François Le Gall)
  - Currently Strassen's is better **in practice**.
  - There is a group theoretic conjecture implying  $\Theta(n^2)$ .

- Implement Strassen's algorithm.