

# The Heap and the Quick

Serge Kruk

October 12, 2021

- Brute-force (sort of)  $O(n^2)$ 
  - Insertion sort
  - Bubble sort
- Divide-and-conquer  $O(n \log n)$ 
  - Merge sort
  - Heap sort
  - Quick sort

# Bubble

```
1 def bs(a=[]):  
2     n = len(a)  
3     for i in range(n-1):  
4         for j in range(i+1,n):  
5             if a[i] > a[j]:  
6                 a[i],a[j] = a[j],a[i]  
7     return a
```

## Loop invariant

Entering the outer loop with  $i$  at value  $k$  the sub-array  $a[0 \dots k-1]$  contains the  $k$  smallest elements of the whole array, in a sorted order.

# Insertion sort

```
1 def isort(a):
2     for i in range(1, len(a)):
3         currentvalue = a[i]
4         position = i
5
6         while position > 0 and a[position-1] > currentvalue:
7             a[position] = a[position-1]
8             position = position-1
9
10        a[position] = currentvalue
11    return a
```


## Loop invariant

Entering the outer loop with  $i$  at value  $k$  the sub-array  $a[0 \dots k]$  is sorted.


# How to merge two sorted arrays

- Input is a pair of sorted arrays  $a$ ,  $b$
- Output is a new array, sorted, containing all elements of  $a$  and  $b$ .


# Merging



# Merging




# Merging







# Merging




# Merging




# Merging




# Merging



# Merging



# Merging



# How to merge two sorted arrays

```
def merge(a,b):  
  
    return c
```

# How to merge two sorted arrays

```
1 ● def merge(a,b):
2     ia,ib,ic,na,nb = 0,0,0,len(a),len(b)
3     nc = na+nb
4     c=[0]*nc
5     while (ic < nc):
6         if (ia < na):
7             if (ib < nb):
8                 if (a[ia] < b[ib]):
9                     c[ic],ic,ia = a[ia],ic+1,ia+1
10                else:
11                    c[ic],ic,ib = b[ib],ic+1,ib+1
12            else:
13                c[ic],ic,ia = a[ia],ic+1,ia+1
14        else:
15            c[ic],ic,ib = b[ib],ic+1,ib+1
16    return c
```



Given

- $n_a$  as the length of a
- $n_b$  as the length of b
- Runtime is  $\Theta(n_a + n_b)$

# Merge sort

General idea:

- Split in half
- Sort the left part
- Sort the right part
- Merge the sorted halves (which we did above)
- We can (and will) call `merge`. Fill-in the missing part.

```
1  def msort(a):  
2      ...  
3      return ...
```

# Merge sort

```
1 def msort(a):
2     n = len(a)
3     m = n//2
4     if n <= 1:
5         return a
6     else:
7         return merge(msort(a[0:m]), msort(a[m:]))
1 print([msort([])==[],
2         msort([4])==[4],
3         msort([1,3,3,2,3,1])==[1,1,2,3,3,3],
4         msort([10,9,8,7,6,5,4,3,2,1]) == [1,2,3,4,5,6,7,8,9,10]])
[True, True, True, True]
```

# Merge sort

- As a one-liner. Just to mess with you.

```
def msort(a):  
    return a if len(a)//2 <= 1 else  
        merge(msort(a[0:len(a)//2]), msort(a[len(a)//2:]))
```

# Merge sort

Correctness proof

By induction

# Proof of correctness

- Base case: if the array is of length 0 or 1, we return this array.
- Induction hypothesis: Assume `msort` works for arrays of length 0 to  $k$ .
  - Consider an array of length any length up to  $k + l$  (as long as  $k + l < 2k$ )
  - We split it in two parts, one of length  $\lfloor (k + l)/2 \rfloor$ , the other  $\lceil (k + l)/2 \rceil$
  - Since both parts are at most length  $k$ , the hypothesis holds and they return sorted.
  - Assuming that `merge` works correctly, we return a sorted array of length  $k + l$ .

## Conclusion

`msort` is correct if `merge` is correct.

# Merge sort runtime

```
def msort(a):  
    m = len(a)//2  
    if m <= 1:  
        return a  
    else:  
        return merge(msort(a[0:m]), msort(a[m:]))
```

Fill-in the recurrence

$$T(n) = \begin{cases} 1 & n \leq 1 \\ ??? & \text{Otherwise} \end{cases}$$

# Merge sort runtime

$$T(n) = \begin{cases} 1 & n \leq 1 \\ 2T(n/2) + n & \text{Otherwise} \end{cases}$$

Solves to  $\Theta(n \log n)$ .



# Trick(?) question

```
def msort(a):  
    m = len(a)//2  
    if m <= 1:  
        return a  
    else:  
        return merge(msort(a[0:m]), msort(a[m:]))
```

If I call `msort([a1, a2, a3, a4, a5, a6, a7, a8])` what are the next three calls to `msort` to start executing?

- `msort([a1, a2, a3, a4, a5, a6, a7, a8])`
  - `msort([a1, a2, a3, a4])`
  - `msort([a1, a2])`
  - `msort([a1])`

# A sort not covered in textbooks anymore

## Shell Sort

For the next few minutes, keep this sequence of integers in mind.

`gaps = [701, 301, 132, 57, 23, 10, 4, 1]`

(It is called a **gap** sequence.)

# Shell sort

```
def shellSort(lst,gaps):  
    for gap in gaps:  
        for i in range(gap):  
            gapISort(lst,i,gap)  
  
def gapISort(lst,start,gap):  
    for i in range(start+gap,len(lst),gap):  
        v = lst[i]  
        p = i  
        while p >= gap and lst[p-gap] > v:  
            lst[p] = lst[p-gap]  
            p = p-gap  
        lst[p]=v
```

# Proof of correctness

The last gap is an insertion sort.

# The gap sequence determines the runtime

- $\lfloor n/2^k \rfloor$  will yield  $\Theta(n^2)$
- $2\lfloor n/2^{k+1} \rfloor + 1$  will yield  $\Theta(n^{\frac{3}{2}})$
- $2^p 3^q$  will yield  $\Theta(n \log^2 n)$
- Better? Best? Is still open!

## Read

- Knuth The Art of Computer Programming for beautiful mathematics.
- On Nostalgia Night, I will tell you why I am fond of ShellSort.

## Runtime

- From  $O(n^2)$  down to  $O(n \log n)$
- Depends on clever procedures

## Now for something completely different

```
def abstractSort(a):  
    n,na = len(a), []  
    for i in range(n):  
        smallest = extractSmallestAndDelete(a)  
        append smallest to na  
    return na
```

- We could easily do this in  $O(n^2)$ .
- We will attempt to do it more efficiently.




# Consider the following data structure properties

- Binary Tree-like, but always balanced (unlike BST)
- Invariant: the value of a node is always less than the value of its children.

It exists.

It is called a min-heap

## Graphical representation



How would you do it?

# Here is the “standard” implementation

- Store in an array of size  $1 + n$

# Here is the “standard” implementation

- Store in an array of size  $1 + n$
- The number of elements is in position 0

# Here is the “standard” implementation

- Store in an array of size  $1 + n$
- The number of elements is in position 0
- The Root is at position 1

# Here is the “standard” implementation

- Store in an array of size  $1 + n$
- The number of elements is in position 0
- The Root is at position 1
- The Left child of node at position  $i$  is at  $2i$

# Here is the “standard” implementation

- Store in an array of size  $1 + n$
- The number of elements is in position 0
- The Root is at position 1
- The Left child of node at position  $i$  is at  $2i$
- The Right child of node at position  $i$  is at  $2i+1$



# Here is the “standard” implementation

- Store in an array of size  $1 + n$
- The number of elements is in position 0
- The Root is at position 1
- The Left child of node at position  $i$  is at  $2i$
- The Right child of node at position  $i$  is at  $2i+1$
- Therefore, the parent of  $i$  is at  $\lfloor i/2 \rfloor$

# Defining property

## Min-heap property

An array  $a[1..n]$  has the min-heap property if

$$a[\lfloor i/2 \rfloor] \leq a[i] \quad 2 \leq i \leq n$$

- Note that this does **not** imply the array is sorted.

# Building the heap

- We need to build this heap, one element at a time.

# Building the heap

- We need to build this heap, one element at a time.
- We will assume that we know ahead of time how large it can grow.

# Building the heap


- We need to build this heap, one element at a time.
- We will assume that we know ahead of time how large it can grow.
- Start, we allocate an array of size one more and put 0 at position 0.

# Building the heap

- We need to build this heap, one element at a time.
- We will assume that we know ahead of time how large it can grow.
- Start, we allocate an array of size one more and put 0 at position 0.
- Adding the first element is simple. You see this?


# How do we add a node? Say value 13

Step 0 : Add it at the end



# How do we add a node? Say value 13


Step 1 : Swap with its parent if necessary





# How do we add a node? Say value 13

Step 2 : Repeat until min-heap property is restored



Let us decompose this into

```
def newheap(n):  
    return [0]*(n+1)  
def insert(a,e):  
    # Inserting element e into min-heap a  
    a[0] = a[0] + 1  
    a[a[0]] = e  
    heapfixup(a,a[0])
```

# Implement this “fixing up”

```
def heapfixup(a,i):  
    # Fix up from position i to restore  
    # min-heap property of heap a
```

# Implementation

```
def heapfixup(a,i):  
    while i > 1:  
        p = i // 2  
        if a[p] > a[i]:  
            a[p],a[i] = a[i],a[p]  
            i = p  
        else:  
            return
```

# Small tests

```
v = [1,4,2,5,3,6,4,7]
h = newheap(len(v))
print(h)
for e in v:
    insert(h,e)
    print(h)
```

```
[0, 0, 0, 0, 0, 0, 0, 0, 0]
[1, 1, 0, 0, 0, 0, 0, 0, 0]
[2, 1, 4, 0, 0, 0, 0, 0, 0]
[3, 1, 4, 2, 0, 0, 0, 0, 0]
[4, 1, 4, 2, 5, 0, 0, 0, 0]
[5, 1, 3, 2, 5, 4, 0, 0, 0]
[6, 1, 3, 2, 5, 4, 6, 0, 0]
[7, 1, 3, 2, 5, 4, 6, 4, 0]
[8, 1, 3, 2, 5, 4, 6, 4, 7]
```

Notice that the heap condition holds at every step

- LI: Descendants of position  $i$  are larger than  $a[i]$ .
- $\Theta(\log n)$  where  $n$  is the size of the heap since we divide by two at every step.

# Recall our abstract sort

```
def abstractSort(a):  
    n, na = len(a), []  
    for i in range(n):  
        smallest= extractSmallestAndDelete(a)  
        na.append(smallest)  
    return na
```

- Where is the smallest? At position 1.
- Now what?

# Extract smallest and delete

How would you do it?



# Extract smallest and delete from heap

- Save the element at position 1
- Move element  $n$  to position 1
- Fix heap property downward

# Implementation

```
def extractsmallest(a):  
    e,a[1],a[0] = a[1],a[a[0]],a[0]-1  
    heapfixdown(a,1)  
    a[a[0]+1]=0  
    return e
```

# Implementation

```
def heapfixdown(a,i):
```

# Implementation

```
def heapfixdown(a,i):
    while 2*i <= a[0]:
        c = 2*i
        if c+1 <= a[0]:
            if a[c+1] < a[c]:
                c = c+1
        if a[i] > a[c]:
            a[i],a[c] = a[c],a[i]
            i = c
    else:
        return
```

# Small tests

```
v = [1,4,2,5,3,6,4,7]
h = newheap(len(v))
for e in v:
    insert(h,e)
print(0, "--", h)
for _ in range(h[0]):
    e = extractsmallest(h)
    h[h[0]+1] = e
    print(e, "--", h)
```

```
0 -- [8, 1, 3, 2, 5, 4, 6, 4, 7]
1 -- [7, 2, 3, 4, 5, 4, 6, 7, 1]
2 -- [6, 3, 4, 4, 5, 7, 6, 2, 1]
3 -- [5, 4, 5, 4, 6, 7, 3, 2, 1]
4 -- [4, 4, 5, 7, 6, 4, 3, 2, 1]
4 -- [3, 5, 6, 7, 4, 4, 3, 2, 1]
5 -- [2, 6, 7, 5, 4, 4, 3, 2, 1]
6 -- [1, 7, 6, 5, 4, 4, 3, 2, 1]
7 -- [0, 7, 6, 5, 4, 4, 3, 2, 1]
```

- LI: Every ancestor of position  $i$  is smaller than  $a[i]$ .
- $\Theta(\log n)$

# We have all we need to create our sort

```
def abstractSort(a):  
    n, na = len(a), []  
    for i in range(n):  
        smallest= extractSmallestAndDelete(a)  
        na.append(smallest)  
    return na
```

# Heapsort

```
def heapsort(x):  
    n = len(x)  
    a = newheap(n)  
    for i in range(n):  
        insert(a,x[i])  
    for i in range(n):  
        x[i] = extractsmallest(a)  
    return x
```



- Runtime :  $\Theta(n \log n)$
- Space required:  $2n$

The distinction between algorithms and data structure is fuzzy!

# Extracting the data structure of heap sort

The operations:

- Insert
- Extract minimum and delete

Form the basis of an abstract data structure called a **priority queue**

Applications include scheduling jobs by an operating system.


# Various implementation of priority queues

	Insert	Extract
Sorted array	$n$	1
Min heap	$\log n$	$\log n$
Unsorted array	1	$n$

# Project: Heapsort can be improved

Hint: Use the space of the array for the heap. Discuss

# Improved heapsort (graphical view)



# Improved heapsort

- Runtime is still  $\Theta(n \log n)$
- Space requirement is  $n$

# Homework/Test questions

- Insert the following items in a min-heap, 18,5,19,3,27,11. Draw the heap at each step.
- Is an array, sorted in non-decreasing order, a min-heap?
- Is a min-heap an array in non-decreasing order?



## Problem

Given an (unsorted) array of  $n$  elements, return the largest  $k$  elements.

- Simplistic, obvious approach? ( $n \log n$ )
- You are asked to do it in ( $n \log k$ )

```
1  def largest_k(a,k=1):  
2      """a is an array of n elements. We return the k largest."""  
3      return b
```

# Sorts up to now

- Brute-force (sort of)  $O(n^2)$ 
  - Insertion sort
  - Bubble sort
- Clever modification of insertion sort
  - Shell  $O(n^{3/2})$
  - Shell  $O(n \log^2 n)$
- Divide-and-conquer  $O(n \log n)$ 
  - Merge sort
- New data structure  $O(n \log n)$ 
  - Heap sort

# Now for the last "Classical" sort

Can you guess?

*Algol 60 is a language so far ahead of its time that it was not only an improvement on its predecessors but also on nearly all of its successors.*

*Due credit must be paid to the designers of Algol 60 who included recursion in their language and allowed me to describe my invention (quicksort) so elegantly to the world.*

# Bubble sort before Algol

```
SUBROUTINE sort (array_x, array_y, datasize)
  REAL array_x(*)
  REAL array_y(*)
  INTEGER datasize
  REAL x_temp
  REAL y_temp
  LOGICAL inorder
  inorder = .false.
```

```
do 90 while (inorder.eq..false.)
  inorder = .true.
do 91 i=1, (datasize-1)
  if (array_x(i).eq.array_x(i+1) ) then
    if (array_y(i).lt.array_y(i+1) ) then
      x_temp = array_x(i)
      y_temp = array_y(i)
      array_x(i) = array_x(i+1)
      array_y(i) = array_y(i+1)
      array_x(i+1) = x_temp
      array_y(i+1) = y_temp
      inorder = .false.
    endif
  endif
endif
```

```
if (array_x(i).lt.array_x(i+1) )then
  x_temp = array_x(i)
  y_temp = array_y(i)
  array_x(i) = array_x(i+1)
  array_y(i) = array_y(i+1)
  array_x(i+1) = x_temp
  array_y(i+1) = y_temp
  inorder = .false.
endif
91  continue
90  continue
END SUBROUTINE sort
```

# Bubble sort in Algol


```
PROC sort = (REF[]DATA array)VOID:
(
  BOOL sorted;
  FOR size FROM UPB array - 1 BY -1 WHILE
    sorted := TRUE;
    FOR i FROM LWB array TO size DO
      IF array[i+1] < array[i] THEN
        swap(array[i:i+1]);
        sorted := FALSE
      FI
    OD;
    NOT sorted
  DO SKIP OD
);
```



- Before algol
  - Fortran (imperative)
  - Lisp (functional)
- Algol introduced
  - Coercion
  - Structures (unions)
  - Local variables
  - Recursion into imperative programs
- First compiler by Dijkstra and Zonneveld
  - Fits in 4K bytes
  - Written in a few months by two mathematicians

# Quicksort (the most used sorting algo)

Illustration of brilliant idea. Say we have this array



Pick one element, say 55 then **partition** by moving smaller than 55 to the left and the larger (or equal) to the right.



Note that 55 is in its final place! (This is **key**)  
Now recurse on both the left and right subarrays.

# Partition

- Pick the first element of the array as the **pivot**.
- Move all smaller elements to the left of the pivot.
- Move all larger elements to the right of the pivot.


```
def partition(a,l,u)
```

```
    return m
```

# Partition

```
def partition(a,l,u):  
    t = a[l]  
    m = l  
    for i in range(l+1,u+1):  
        if a[i] < t:  
            m = m+1  
            a[i],a[m] = a[m],a[i]  
    a[m],a[l] = a[l],a[m]  
    return m
```

Loop invariant



# Implement Quicksort (recursively!)

- Hint: Divide and conquer
- Assume you have `partition`

```
def qsort(a):
```

```
    return a
```

# Implementation

```
def qsort0(a,l=0,u=None):  
    if u is None:  
        u = len(a)-1  
    if l < u:  
        m = partition(a,l,u)  
        qsort0(a,l,m-1)  
        qsort0(a,m+1,u)  
    return a
```

- Assume that we will partition in  $O(n)$ .
- Assume that we partition equally.

$$T(n) = \begin{cases} C & n \leq n_0 \\ 2T(n/2) + n & \text{otherwise} \end{cases}$$

$$\begin{aligned}T(n) &= 2 T(n/2) + n \\&= 2[2 T(n/4) + n/2] + n &= 4 T(n/4) + 2n \\&= 4[2 T(n/8) + n/4] + 2n &= 8 T(n/8) + 3n \\&= \vdots \\&= 2^k T(n/2^k) + kn\end{aligned}$$

Since we do  $\log n$  steps, we obtain  $T(n) = \Theta(n \log n)$



# A dare

For those of you who are still reluctant to consider recursion, I dare you to write quicksort without recursion (and get it right). It's doable, of course, but it's a nightmare to get right.

- What if we partition with only **one** element on one side?

$$\begin{aligned}T(n) &= T(n-1) + T(1) + n \\&= T(n-2) + T(1) + n-1 + T(1) + n \\&= \dots \\&= T(1) + T(1) + \dots + T(1) + 1 + 2 \\&\quad + \dots + n-1 + n \\&\in \Theta(n^2)\end{aligned}$$

# Pivot choice

- 1 What happens if the partition is bad, say 10%, 90%?
- 2 Randomly is a possibility
- 3 Some form of median is best

Read the paper by McIlroy

# Better implementation?

```
def qsort1(a,l=0,u=None):  
    if u is None:  
        u = len(a)-1  
    if l < u-128:  
        m = partition(a,l,u)  
        qsort1(a,l,m-1)  
        qsort1(a,m+1,u)  
    else l < u:  
        insertion_sort(a,l,u)  
    return a
```

Why?

C

```
void qsort(void *base,  
           size_t nmemb,  
           size_t size,  
           int (*compar)(const void *, const void *));
```

## Java

```
public <A extends Comparable<? super A>>  
    void sort(List<A> list) { }
```

## Go

```
func Sort(data Interface)
type Interface interface {
    Len() int
    Less(i, j int) bool
    Swap(i, j int)
}
```



- Contrast
  - Textbook (Knuth):  $\text{overhead} \approx \text{comparisons} < \text{swaps}$
  - Findings:  $\text{overhead} < \text{swaps} < \text{comparisons}$
- Why? Because `cmp` is a function call (interpreted).
- Textbook presentation was flawed (from 1970 to 2000).
- Which explains why we are counting comparisons.

# Homework/Test questions

- Why does Quicksort switch to insertion sort?
- What happens if the partition is 90 - 10?
- Which pivot choice is used by Bentley and McIlroy?
- Why would you not use a random pivot?
- What if there are multiple repeated elements? (Better partition?)
- Heapsort and Quicksort are both  $n \log n$ . Tradeoffs?
- Why would one try an iterative implementation of either?
- Which sort would be easier to implement in parallel?

# Class project

You have two datasets indexed by student ID, one large and one small. Your task is to extract the elements that appear in both sets? You can call a sorting routine (no need to code it). To simplify let us assume that the input to your function is a pair of arrays. In each array, the elements are a tuple

Student ID	(9digits)
------------	-----------

Student record (thousands of bytes)
-------------------------------------

```
def extract_common(A, B):  
    """ A is HUGE; B is relatively small """  
    return C
```

Try to do this as efficiently as possible.